

Breast Cancer Mini Project

Gavin Ambrose A18548522

Table of contents

Background	1
Data Import	1
Principal Component Analysis	4
Interpreting PCA results	6
Variance explained	8
Communicating PCA results	10
Hierarchical clustering	11
Selecting number of clusters	12
Using different methods	13
Combining methods	14
Sensitivity/Specificity	15
Prediction	16

Background

In today's class, we will apply the methods and techniques of clustering and PCA to help make sense of a real world breast cancer fine needle aspiration (FNA) biopsy data set.

Data Import

We start by importing our data. It is a CSV file, so we will use the `read.csv` function.

```
fna.data <- "WisconsinCancer.csv"  
  
wisc.df <- read.csv(fna.data, row.names = 1)  
head(wisc.df, 4)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419		0.07871	1.0950	0.9053
842517	0.1812		0.05667	0.5435	0.7339
84300903	0.2069		0.05999	0.7456	0.7869
84348301	0.2597		0.09744	0.4956	1.1560
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003		0.006193	25.38	17.33
842517	0.01389		0.003532	24.99	23.41
84300903	0.02250		0.004571	23.57	25.53
84348301	0.05963		0.009208	14.91	26.50
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.60	2019.0	0.1622	0.6656	
842517	158.80	1956.0	0.1238	0.1866	
84300903	152.50	1709.0	0.1444	0.4245	
84348301	98.87	567.7	0.2098	0.8663	
	concavity_worst	concave.points_worst	symmetry_worst		
842302	0.7119		0.2654	0.4601	
842517	0.2416		0.1860	0.2750	
84300903	0.4504		0.2430	0.3613	
84348301	0.6869		0.2575	0.6638	
	fractal_dimension_worst				
842302		0.11890			
842517		0.08902			
84300903		0.08758			
84348301		0.17300			

```
# Removing the pathologist diagnosis so that we can do it ourselves  
wisc.data <- wisc.df[,-1]  
  
#Setting up a diagnosis vector for later  
diagnosis <- wisc.df$diagnosis
```

Q1. How many observations are in this dataset?

```
dim(wisc.data)
```

```
[1] 569 30
```

There are 569 observations

Q2. How many of the observations have a malignant diagnosis?

```
table(diagnosis)
```

```
diagnosis  
B M  
357 212
```

There are 212 observations that have a malignant diagnosis

Q3. How many variables/features in the data are suffixed with `_mean`?

```
length(grep("_mean", colnames(wisc.data)))
```

```
[1] 10
```

There are 10 variables/features in the data suffixed with “`_mean`”

Principal Component Analysis

The next step in your analysis is to perform principal component analysis (PCA) on wisc.data.

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03

symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

```
wisc.pr <- prcomp(wisc.data, scale = T)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)?

44.2% of the original data is captured by PCA1.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

Three PCAs are needed to record the original caruance if the data.


Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

To record 90% of the original variance, 7 PCAs are needed.

Interpreting PCA results

Now you will use some visualizations to better understand your PCA model. A common visualization for PCA results is the so-called biplot.

```
biplot(wisc.pr)
```



Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?


This graph is not easy to understand, there is a lot of overlapping text in similar colors that makes individual components hard to determine.

Let's Use ggplot instead

```
library(ggplot2)
```


```
Warning: package 'ggplot2' was built under R version 4.4.3
```

```
ggplot(wisc.pr$x) +  
  aes(PC1, PC2, col=diagnosis) +  
  geom_point()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
library(ggplot2)  
  
ggplot(wisc.pr$x) +  
  aes(PC1, PC3, col=diagnosis) +  
  geom_point()
```



The variance displayed in the points is less in the second plot. The points are scattered more in the first graph. Another difference is that the y axis values change, decreasing in the secind plot

Variance explained

We will now use a scree plot to show how much variance each PC captures.

We will calculate variance by squaring the `sdev` component of `wisc.pr`

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

We will now calculate the variance of each principle component by dividing by the variance of all principle components.

```
pve <- pr.var / sum( pr.var )
plot(c(1,pve), xlav = "Principal Component", ylab = "Proportion of Variance Explained", ylim
```

```

Warning in plot.window(...): "xlav" is not a graphical parameter


Warning in plot.xy(xy, type, ...): "xlav" is not a graphical parameter

Warning in axis(side = side, at = at, labels = labels, ...): "xlav" is not a
graphical parameter
Warning in axis(side = side, at = at, labels = labels, ...): "xlav" is not a
graphical parameter

Warning in box(...): "xlav" is not a graphical parameter

Warning in title(...): "xlav" is not a graphical parameter

```




Let's create an alternate scree plot of the data using `barplot()`

```

barplot(pve, ylab = "Percent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )

```



Communicating PCA results

In this section we will check your understanding of the PCA results, in particular the “loadings” and “variance explained”.

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

```
wisc.pr$rotation[,1]
```

radius_mean	texture_mean	perimeter_mean
-0.21890244	-0.10372458	-0.22753729
area_mean	smoothness_mean	compactness_mean
-0.22099499	-0.14258969	-0.23928535
concavity_mean	concave.points_mean	symmetry_mean
-0.25840048	-0.26085376	-0.13816696
fractal_dimension_mean	radius_se	texture_se
-0.06436335	-0.20597878	-0.01742803
perimeter_se	area_se	smoothness_se
-0.21132592	-0.20286964	-0.01453145

compactness_se	concavity_se	concave.points_se
-0.17039345	-0.15358979	-0.18341740
symmetry_se	fractal_dimension_se	radius_worst
-0.04249842	-0.10256832	-0.22799663
texture_worst	perimeter_worst	area_worst
-0.10446933	-0.23663968	-0.22487053
smoothness_worst	compactness_worst	concavity_worst
-0.12795256	-0.21009588	-0.22876753
concave.points_worst	symmetry_worst	fractal_dimension_worst
-0.25088597	-0.12290456	-0.13178394

The `concave.points_mean` principle component is -0.26085376. There are not any components with larger contributions. The negative number represents how the feature varies relative to the PC axis.

Hierarchical clustering

Q.10 Using the `plot` and `abline()` functions, what is the height at which the clustering model has 4 clusters?


```
data.scaled <- scale(wisc.data)

data.dist <- dist(data.scaled)

wisc.hclust <- hclust(data.dist, method = "complete")

plot(wisc.hclust)
abline(h=19, col = "red", lty=2)
```

Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

The height necessary to create a model with 4 clusters is 19. Although there are other possible heights that would meet the same outcome, 19 was chosen

Selecting number of clusters

```
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)  
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

Q11. OPTIONAL: Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 6? How do you judge the quality of your result in each case?

I could not find a better cluster vs diagnosis match by changing the number of clusters. There was not a greater separation between malignant and benign that I could determine using clusters between 2 and 6.


Using different methods

There are number of different “methods” we can use to combine points during the hierarchical clustering procedure. These include `single`, `complete`, `average` and (my favorite) `ward.D2`.

Q12. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning.

```
wisc.hclustdiffmethod <- hclust(data.dist, method = "ward.D2")

plot(wisc.hclustdiffmethod)
abline(h=60, col = "red", lty=2)
```



data.dist
hclust (*, "ward.D2")

```
wisc.hclust.clustersdiffmethod <- cutree(wisc.hclustdiffmethod, k = 4)
table(wisc.hclust.clustersdiffmethod, diagnosis)
```


diagnosis		
wisc.hclust.clustersdiffmethod	B	M
1	0	115
2	6	48
3	337	48
4	14	1

I like the `ward.D2` because it gives much more discernable “goalposts”, and the clustering with $k = 4$ presented a much more distinguishable difference between benign and malignant samples.

Combining methods

We have tried PCA and hierarchical clustering separately. Now let's combine them: cluster on the PC scores instead of the original 30 features.

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:7]), method = "ward.D2")
plot(wisc.pr.hclust)
```



```
dist(wisc.pr$x[, 1:7])
hclust (*, "ward.D2")
```

Note the two main branches of our dendrogram indicating two main clusters - maybe these are malignant and benign. Let's find out!

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
 1 2
216 353
```

I want to know how the clusterinf in grps with values 1 or 2 correspond to the expert diagnosis


```
table(grps, diagnosis)
```

grps	B	M
1	28	188
2	329	24

Sensitivity/Specificity

FP: 28 TP: 188 TN: 329 FN: 24 Sensitivity: $TP/(TP + FN)$ Specificity: $TN/(TN+FP)$

```
ggplot(wisc.pr$x) +
  aes(PC1, PC2, color = factor(grps)) +
  geom_point(size = 1) +
  scale_color_manual(
    values = c("1" = "blue", "2" = "red"), labels = c("Benign", "Malignant")) +
  labs(color = "Diagnosis")
```



Q13. How well does the newly created hclust model with two clusters separate out the two “M” and “B” diagnoses?

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)

table(wisc.pr.hclust.clusters, diagnosis)
```

diagnosis		
wisc.pr.hclust.clusters	B	M
1	28	188
2	329	24

There is no significant difference between the two clusters model and the 4 cluster model

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.hclust.clusters and wisc.pr.hclust.clusters) with the vector containing the actual diagnoses.

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)

table(wisc.hclust.clusters, diagnosis)
```

diagnosis		
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

Only the ward method separate out the data properly with 2 clusters, complete also worked but it needed 4 clusters

Prediction

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

```


          PC1        PC2        PC3        PC4        PC5        PC6        PC7
[1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
[2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
          PC8        PC9        PC10       PC11       PC12       PC13       PC14
[1,] -0.2307350  0.1029569 -0.9272861  0.3411457  0.375921  0.1610764  1.187882
[2,] -0.3307423  0.5281896 -0.4855301  0.7173233 -1.185917  0.5893856  0.303029
          PC15       PC16       PC17       PC18       PC19       PC20
[1,]  0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
[2,]  0.1299153  0.1448061 -0.40509706  0.06565549  0.25591230 -0.4289500
          PC21       PC22       PC23       PC24       PC25       PC26
[1,]  0.1228233  0.09358453  0.08347651  0.1223396  0.02124121  0.078884581
[2,] -0.1224776  0.01732146  0.06316631 -0.2338618 -0.20755948 -0.009833238
          PC27       PC28       PC29       PC30
[1,]  0.220199544 -0.02946023 -0.015620933  0.005269029
[2,] -0.001134152  0.09638361  0.002795349 -0.019015820

```

```

plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")

```



Q16. Which of these new patients should we prioritize for follow up based on your results?

I would prioritize treatments for patient 1 because they are more closely seen in the malignant side.