

Class9Candy_Project

Gavin Ambrose PID A18548522

Table of contents

Importing candy data	1
What is in the dataset	2
What is your favorite candy?	3
Exploratory analysis	5
Overall Candy Rankings	6
Time to add some useful color	10
Taking a look at pricepercent	11
Make a lollipop chart of pricepercent	13
Exploring the correlation structure	14
Principal Component Analysis	15
Summary	20

Importing candy data

First things first, let's get the data from the FiveThirtyEight GitHub repo

```
candy_file <- read.csv("https://raw.githubusercontent.com/fivethirtyeight/data/master/candy-  
head(candy_file)
```

	chocolate	fruity	caramel	peanutyalmondny	nougat	crispedricewafer	
100 Grand	1	0	1		0	0	1
3 Musketeers	1	0	0		0	1	0
One dime	0	0	0		0	0	0
One quarter	0	0	0		0	0	0
Air Heads	0	1	0		0	0	0
Almond Joy	1	0	0		1	0	0
	hard	bar	pluribus	sugarpercent	pricepercent	winpercent	

100 Grand	0	1	0	0.732	0.860	66.97173
3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

What is in the dataset

The dataset includes all sorts of information about different kinds of candy. For example, is a candy chocolaty? Does it have nougat? How does its cost compare to other candies? How many people prefer one candy over another?

Q1. How many different candy types are in this dataset?

```
str(candy_file)
```

```
'data.frame': 85 obs. of 12 variables:
 $ chocolate      : int  1 1 0 0 0 1 1 0 0 0 ...
 $ fruity          : int  0 0 0 0 1 0 0 0 0 1 ...
 $ caramel         : int  1 0 0 0 0 0 1 0 0 1 ...
 $ peanutyalmondy : int  0 0 0 0 0 1 1 1 0 0 ...
 $ nougat          : int  0 1 0 0 0 0 1 0 0 0 ...
 $ crispedricewafer: int  1 0 0 0 0 0 0 0 0 0 ...
 $ hard             : int  0 0 0 0 0 0 0 0 0 0 ...
 $ bar              : int  1 1 0 0 0 1 1 0 0 0 ...
 $ pluribus         : int  0 0 0 0 0 0 0 1 1 0 ...
 $ sugarpercent    : num  0.732 0.604 0.011 0.011 0.906 ...
 $ pricepercent    : num  0.86 0.511 0.116 0.511 0.511 ...
 $ winpercent      : num  67 67.6 32.3 46.1 52.3 ...
```

There are 85 different candy types within this dataset

Q2. How many fruity candy types are in the dataset?

```
sum(candy_file$fruity)
```

[1] 38

There are 38 fruity candy types in the dataset

What is your favorite candy?

One of the most interesting variables in the dataset is `winpercent`. For a given candy this value is the percentage of people who prefer this candy over another randomly chosen candy from the dataset. Higher values indicate a more popular candy.

Q3. What is your favorite candy (other than Twix) in the dataset and what is it's `winpercent` value?

```
candy_file["Mike & Ike",]$winpercent
```

```
[1] 46.41172
```

My favorite candy is Mike and Ike's and they have a `winpercent` value of 46.411

Q4. What is the `winpercent` value for "Kit Kat"?

```
candy_file["Kit Kat",]$winpercent
```

```
[1] 76.7686
```

Kit Kat's have a `winpercent` of 76.768

Q5. What is the `winpercent` value for "Tootsie Roll Snack Bars"?

```
candy_file["Tootsie Roll Snack Bars",]$winpercent
```

```
[1] 49.6535
```

Tootsie Roll Snack Bars have a `winpercent` of 49.653

```
library("skimr")
```

```
Warning: package 'skimr' was built under R version 4.4.3
```

```
skim(candy_file)
```

Table 1: Data summary

Name	candy_file
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	com-	mean	sd	p0	p25	p50	p75	p100	hist
		plete_rate								
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyal-	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
mondy										
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedrice-	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
wafer										
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

The winpercent column is on a different scale than the majority of the data because it has a mean data of 50 whereas the others are between 0.08 and 0.4 means.

Q7. What do you think a zero and one represent for the candy\$chocolate column?

I think that a one represents a belonging to that category wheras a 0 represents a lack of belonging to that category

Exploratory analysis

A good place to start any exploratory analysis is with a histogram.


Q8. Plot a histogram of winpercent values

```
library(ggplot2)
```

```
Warning: package 'ggplot2' was built under R version 4.4.3
```

```
ggplot(candy_file) + aes(winpercent) + geom_histogram() + stat_bin (bins = 10)
```

```
`stat_bin()` using `bins = 30`. Pick better value `binwidth`.
```



Q9. Is the distribution of winpercent values symmetrical?

The histogram is slightly skewed to the right. Making it asymmetric

Q10. Is the center of the distribution above or below 50%?

The center of the distribution is below 50%

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
mean_chocolate <- mean(candy_file$winpercent[as.logical(candy_file$chocolate)])
mean_chocolate
```

```
[1] 60.92153
```

```
mean_fruity <- mean(candy_file$winpercent[as.logical(candy_file$fruity)])
mean_fruity
```

```
[1] 44.11974
```

Chocolate candy has a higher rating than fruity on average

Q12. Is this difference statistically significant?

```
t.test(candy_file$winpercent[as.logical(candy_file$chocolate)], candy_file$winpercent[as.lo
```

Welch Two Sample t-test

```
data: candy_file$winpercent[as.logical(candy_file$chocolate)] and candy_file$winpercent[as.lo
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974
```

We can reject the null hypothesis that there is a difference. The difference is significant because a $p=2.871e-08 < 0.05$

Overall Candy Rankings

Let's use the base R `order()` function together with `head()` to sort the whole dataset by `winpercent`. Or if you have been getting into the tidyverse and the `dplyr` package you can use the `arrange()` function together with `head()` to do the same thing and answer the following questions:

Q13. What are the five least liked candy types in this set?

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
candy_file |> arrange(winpercent) |> head(5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat	
Nik L Nip	0	1	0		0	0	
Boston Baked Beans	0	0	0		1	0	
Chiclets	0	1	0		0	0	
Super Bubble	0	1	0		0	0	
Jawbusters	0	1	0		0	0	
	crispedrice	wafers	hard	bar	pluribus	sugarpercent	pricepercent
Nik L Nip	0	0	0		1	0.197	0.976
Boston Baked Beans	0	0	0		1	0.313	0.511
Chiclets	0	0	0		1	0.046	0.325
Super Bubble	0	0	0		0	0.162	0.116
Jawbusters	0	1	0		1	0.093	0.511
	winpercent						
Nik L Nip	22.44534						
Boston Baked Beans	23.41782						
Chiclets	24.52499						
Super Bubble	27.30386						
Jawbusters	28.12744						

The five least liked candies are Nik L Nip, Boston Baked Beans, Chiclets, Super Bubble, and Jawbusters.

Q14. What are the top 5 all time favorite candy types out of this set?


```
candy_file |> arrange(winpercent) |> tail(5)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat
Snickers	1	0	1		1	1
Kit Kat	1	0	0		0	0
Twix	1	0	1		0	0
Reese's Miniatures	1	0	0		1	0
Reese's Peanut Butter cup	1	0	0		1	0
	crispedrice	wafers	hard bar	pluribus	sugar	percent
Snickers	0	0	1		0	0.546
Kit Kat	1	0	1		0	0.313
Twix	1	0	1		0	0.546
Reese's Miniatures	0	0	0		0	0.034
Reese's Peanut Butter cup	0	0	0		0	0.720
	price	percent	winpercent			
Snickers	0.651	76.67378				
Kit Kat	0.511	76.76860				
Twix	0.906	81.64291				
Reese's Miniatures	0.279	81.86626				
Reese's Peanut Butter cup	0.651	84.18029				

The top 5 of all time are Snickers, Kit Kat, Reese's Miniatures, and Reese's Peanut Butter cup.


Q15. Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy_file) + aes(winpercent, rownames(candy_file)) + geom_col()
```



Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent?

```
ggplot(candy_file) + aes(winpercent, reorder(rownames(candy_file), winpercent)) + geom_col()
```




Time to add some useful color

Let's setup a color vector (that signifies candy type) that we can then use for some future plots. We start by making a vector of all black values (one for each candy). Then we overwrite chocolate (for chocolate candy), brown (for candy bars) and red (for fruity candy) values.

```
my_cols=rep("black", nrow(candy_file))
my_cols[as.logical(candy_file$chocolate)] = "chocolate"
my_cols[as.logical(candy_file$bar)] = "brown"
my_cols[as.logical(candy_file$fruity)] = "pink"
```

```
ggplot(candy_file) + aes(winpercent, reorder(rownames(candy_file), winpercent)) + geom_col(f
```



Q17. What is the worst ranked chocolate candy?

Sixlets

Q18. What is the best ranked fruity candy?

Starbursts

Taking a look at pricepercent


What about value for money? What is the best candy for the least money?

```
library(ggrepel)
```

Warning: package 'ggrepel' was built under R version 4.4.3

```
# How about a plot of win vs price
ggplot(candy_file) +
  aes(winpercent, pricepercent, label=rownames(candy_file)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 5)
```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

```
ordleast <- order(candy_file$pricepercent, decreasing = TRUE)
head( candy_file[ordleast,c(11,12)], n=5 )
```

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050


Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
ordmost <- order(candy_file$pricepercent, decreasing = F)
head( candy_file[ordmost,c(11,12)], n=5 )
```

	pricepercent	winpercent
Tootsie Roll Midgies	0.011	45.73675
Pixie Sticks	0.023	37.72234
Dum Dums	0.034	39.46056
Fruit Chews	0.034	43.08892
Strawberry bon bons	0.058	34.57899


Q21. Make a barplot again with `geom_col()` this time using `pricepercent` and then improve this step by step, first ordering the x-axis by value and finally making a so called “dot chat” or “lollipop” chart by swapping `geom_col()` for `geom_point()` + `geom_segment()`.

```
ggplot(candy_file) + aes(pricepercent, reorder(rownames(candy_file), pricepercent)) + geom_c
```



Make a lollipop chart of pricepercent

```
ggplot(candy_file) +
  aes(pricepercent, reorder(rownames(candy_file), pricepercent)) +
  geom_segment(aes(yend = reorder(rownames(candy_file), pricepercent),
                   xend = 0), col="gray40") +
  geom_point()
```



One of the most interesting aspects of this chart is that a lot of the candies share the same ranking, so it looks like quite a few of them are the same price.

Exploring the correlation structure

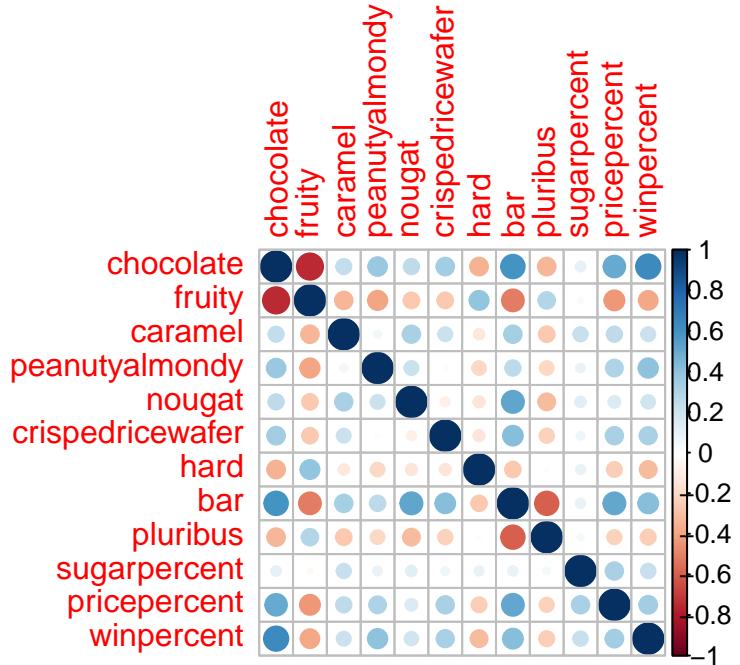
Now that we've explored the dataset a little, we'll see how the variables interact with one another. We'll use correlation and view the results with the corrplot package to plot a correlation matrix.

```
library(corrplot)
```

```
Warning: package 'corrplot' was built under R version 4.4.3
```

```
corrplot 0.95 loaded
```

```
corcandy <- cor(candy_file)
corrplot(corcandy)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

The two variables that have a strong negative correlation chocolate and fruity

Q23. Similarly, what two variables are most positively correlated?

The chocolate variable and the winpercent variable have the strongest positive correlation.

Principal Component Analysis

Let's apply PCA using the prcomp() function to our candy dataset remembering to set the scale=TRUE argument.

```
pcacandy <- prcomp(candy_file, scale = T)
summary(pcacandy)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		

```
Standard deviation      0.74530 0.67824 0.62349 0.43974 0.39760
Proportion of Variance 0.04629 0.03833 0.03239 0.01611 0.01317
Cumulative Proportion  0.89998 0.93832 0.97071 0.98683 1.00000
```


Let's plot our main PCA scores (PCA1 and PCA2)

```
plot(pcacandy$x[,1:2])
```




We can change the colors and characters

```
plot(pcacandy$x[,1:2], col = my_cols, pch = 16)
```



Let's try it with ggplot

```
my_data <- cbind(candy_file, pcacandy$x[,1:3])
pcaggplot <- ggplot(my_data) +
  aes(x=PC1, y=PC2,
      size=winpercent/100,
      text=rownames(my_data),
      label=rownames(my_data)) +
  geom_point(col=my_cols)
pcaggplot
```




Lets use greppel and label the plot

```
pcagggplot + geom_text_repel(size=3.3, col=my_cols, max.overlaps = 7) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
       subtitle="Colored by type: chocolate bar (dark brown), chocolate other (light brown),
       caption="Data from 538")
```

Warning: ggrepel: 59 unlabeled data points (too many overlaps). Consider increasing max.overlaps

Halloween Candy PCA Space


Colored by type: chocolate bar (dark brown), chocolate other (light brown),



Data from 538

Q24. Complete the code to generate the loadings plot above. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you? Where did you see this relationship highlighted previously?

```
ggplot(pcacandy$rotation) +
  aes(x = PC1, y = reorder(rownames(pcacandy$rotation), PC1)) +
  geom_col()
```



The fruity, pluribis, and hard variables are picked up strongly by the PCA. These make sense because they are popular characteristics of candy. This was highlighted in the previous question that also use bar plots.

Summary

Q25. Based on your exploratory analysis, correlation findings, and PCA results, what combination of characteristics appears to make a “winning” candy? How do these different analyses (visualization, correlation, PCA) support or complement each other in reaching this conclusion?

Based on visual data, chocolate and bar candies seems to preform the best, and fruity candies the worst. Positive correlation is associated with chocolate and bar candies, while a negative correlation is associated with fruity candies. PC1 has a shows chocolate and bar variables along one of the axis and fruity and plurabis along a different axis. These all propose that the “winning” candy is likely a chocolate candy in a bar form.