

# Class 14: DESeq2 mini project

Gavin Ambrose PID: A18548522

## Table of contents

(1) Differential Expression Analysis . . . . .	1
Running DESeq2 . . . . .	3
Volcano Plot . . . . .	4
Adding gene annotation . . . . .	6
(2) Pathway Analysis . . . . .	8
(3) Gene Ontology (GO) . . . . .	18
(5) Reactome Analysis . . . . .	20

## (1) Differential Expression Analysis

We need to download both the count and associated metadata

```
library(DESeq2)
```

```
Warning: package 'matrixStats' was built under R version 4.4.3
```

```
metaFile <- "GSE37704_metadata.csv"
countFile <- "GSE37704_featurecounts.csv"

# Import metadata and take a peek
colData = read.csv(metaFile, row.names=1)
head(colData)
```

```
            condition
SRR493366 control_sirna
SRR493367 control_sirna
SRR493368 control_sirna
```

```

SRR493369      hoxa1_kd
SRR493370      hoxa1_kd
SRR493371      hoxa1_kd

# Import countdata
countDataA = read.csv(countFile, row.names=1)
head(countDataA)

      length SRR493366 SRR493367 SRR493368 SRR493369 SRR493370
ENSG00000186092    918      0      0      0      0      0
ENSG00000279928    718      0      0      0      0      0
ENSG00000279457   1982     23     28     29     29     28
ENSG00000278566    939      0      0      0      0      0
ENSG00000273547    939      0      0      0      0      0
ENSG00000187634   3214    124    123    205    207    212
                  SRR493371
ENSG00000186092      0
ENSG00000279928      0
ENSG00000279457     46
ENSG00000278566      0
ENSG00000273547      0
ENSG00000187634    258

```

Now, we must remove the odd first \$length col from countData to get a 1:1 correspondance.

Q. Complete the code below to remove the troublesome first column from countData

```

countData <- as.matrix(countDataA[,-1])
head(countData)

```

```

      SRR493366 SRR493367 SRR493368 SRR493369 SRR493370 SRR493371
ENSG00000186092      0      0      0      0      0      0
ENSG00000279928      0      0      0      0      0      0
ENSG00000279457     23     28     29     29     28     46
ENSG00000278566      0      0      0      0      0      0
ENSG00000273547      0      0      0      0      0      0
ENSG00000187634    124    123    205    207    212    258

```

This looks better but there are lots of zero entries in there so let's get rid of them as we have no data for these.

Q. Complete the code below to filter countData to exclude genes (i.e. rows) where we have 0 read count across all samples (i.e. columns).

```
countData <- countData[rowSums(countData) != 0, ]  
head(countData)
```

	SRR493366	SRR493367	SRR493368	SRR493369	SRR493370	SRR493371
ENSG00000279457	23	28	29	29	28	46
ENSG00000187634	124	123	205	207	212	258
ENSG00000188976	1637	1831	2383	1226	1326	1504
ENSG00000187961	120	153	180	236	255	357
ENSG00000187583	24	48	65	44	48	64
ENSG00000187642	4	9	16	14	16	16

Nice now lets setup the DESeqDataSet object required for the DESeq() function and then run the DESeq pipeline. This is again similar to our last day's hands-on session.

## Running DESeq2

```
dds <- DESeqDataSetFromMatrix(countData= countData,  
                                colData= colData,  
                                design=~condition)
```

Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in design formula are characters, converting to factors

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

mean-dispersion relationship

final dispersion estimates

fitting model and testing

```
dds
```

```
class: DESeqDataSet
dim: 15975 6
metadata(1): version
assays(4): counts mu H cooks
rownames(15975): ENSG00000279457 ENSG00000187634 ...
ENSG00000276345
ENSG00000271254
rowData names(22): baseMean baseVar ... deviance maxCooks
colnames(6): SRR493366 SRR493367 ... SRR493370 SRR493371
colData names(2): condition sizeFactor
```

Next, get results for the HoxA1 knockdown versus control siRNA (remember that these were labeled as “`hoxa1_kd`” and “`control_siRNA`” in our original `colData` metaFile input to `DESeq`, you can check this above and by running `resultsNames(dds)` command).

Q. Call the `summary()` function on your results to get a sense of how many genes are up or down-regulated at the default 0.1 p-value cutoff.

```
res <- results(dds)
summary(res)
```

```
out of 15975 with nonzero total read count
adjusted p-value < 0.1
LFC > 0 (up)      : 4349, 27%
LFC < 0 (down)    : 4396, 28%
outliers [1]       : 0, 0%
low counts [2]     : 1237, 7.7%
(mean count < 0)
[1] see 'cooksCutoff' argument of ?results
[2] see 'independentFiltering' argument of ?results
```

Now we will make a volcano plot, a commonly produced visualization from this type of data that we introduced last day. Basically it's a plot of log2 fold change vs -log adjusted p-value.


## Volcano Plot

```
library(ggplot2)
```

Warning: package 'ggplot2' was built under R version 4.4.3

```
ggplot(res) +  
  aes(res$log2FoldChange,  
      -log(res$padj)) +  
  geom_point()
```

Warning: Removed 1237 rows containing missing values or values outside the scale range  
(`geom\_point()`).



Q. Improve this plot by completing the below code, which adds color, axis labels and cutoff lines:

```
# Make a color vector for all genes  
mycols <- rep("gray", nrow(res) )  
  
# Color blue the genes with fold change above 2  
mycols[ abs(res$log2FoldChange) > 2 ] <- "blue"
```


```

# Color gray those with adjusted p-value more than 0.01
mycols[ res$padj > 0.05 ] <- "gray"

ggplot(res) +
  aes( res$log2FoldChange,
       -log(res$padj)) +
  geom_point(col = mycols) +
  xlab("Log2(FoldChange)") +
  ylab("-Log(P-value)") +
  geom_vline(xintercept = c(-2,2), col = "red", lty = 4) +
  geom_hline(yintercept = -log(0.05), col = "red", lty = 4)

```

Warning: Removed 1237 rows containing missing values or values outside the scale range (`geom\_point()`).



### Adding gene annotation

Since we mapped and counted against the Ensembl annotation, our results only have information about Ensembl gene IDs. However, our pathway analysis downstream will use KEGG pathways, and genes in KEGG pathways are annotated with Entrez gene IDs. So lets add them as we did the last day.

Q. Use the mapIDs() function multiple times to add SYMBOL, ENTREZID and GENENAME annotation to our results by completing the code below.

```
library("AnnotationDbi")
library("org.Hs.eg.db")

columns(org.Hs.eg.db)

[1] "ACNUM"          "ALIAS"           "ENSEMBL"         "ENSEMLPROT"      "ENSEMLTRANS"
[6] "ENTREZID"        "ENZYME"          "EVIDENCE"        "EVIDENCEALL"    "GENENAME"
[11] "GENETYPE"        "GO"               "GOALL"           "IPI"              "MAP"
[16] "OMIM"            "ONTOLOGY"        "ONTOLOGYALL"    "PATH"             "PFAM"
[21] "PMID"            "PROSITE"          "REFSEQ"          "SYMBOL"          "UCSCKG"
[26] "UNIPROT"

res$symbol = mapIds(org.Hs.eg.db,
                     keys=row.names(res),
                     keytype="ENSEMBL",
                     column="SYMBOL",
                     multiVals="first")

'select()' returned 1:many mapping between keys and columns

res$entrez = mapIds(org.Hs.eg.db,
                     keys=row.names(res),
                     keytype="ENSEMBL",
                     column="ENTREZID",
                     multiVals="first")

'select()' returned 1:many mapping between keys and columns

res$name =   mapIds(org.Hs.eg.db,
                     keys=row.names(res),
                     keytype="ENSEMBL",
                     column="GENENAME",
                     multiVals="first")

'select()' returned 1:many mapping between keys and columns
```

```
head(res, 10)
```

```
log2 fold change (MLE): condition hoxa1 kd vs control sirna
Wald test p-value: condition hoxa1 kd vs control sirna
DataFrame with 10 rows and 9 columns
  baseMean log2FoldChange    lfcSE      stat     pvalue
  <numeric>    <numeric> <numeric>    <numeric>    <numeric>
ENSG00000279457  29.913579  0.1792571  0.3248216  0.551863 5.81042e-01
ENSG00000187634 183.229650  0.4264571  0.1402658  3.040350 2.36304e-03
ENSG00000188976 1651.188076 -0.6927205  0.0548465 -12.630158 1.43990e-36
ENSG00000187961 209.637938  0.7297556  0.1318599  5.534326 3.12428e-08
ENSG00000187583 47.255123   0.0405765  0.2718928  0.149237 8.81366e-01
ENSG00000187642 11.979750  0.5428105  0.5215598  1.040744 2.97994e-01
ENSG00000188290 108.922128  2.0570638  0.1969053  10.446970 1.51282e-25
ENSG00000187608 350.716868  0.2573837  0.1027266  2.505522 1.22271e-02
ENSG00000188157 9128.439422 0.3899088  0.0467163  8.346304 7.04321e-17
ENSG00000237330  0.158192   0.7859552  4.0804729  0.192614 8.47261e-01
  padj      symbol      entrez           name
  <numeric> <character> <character> <character>
ENSG00000279457 6.86555e-01       NA        NA          NA
ENSG00000187634 5.15718e-03      SAMD11    148398 sterile alpha motif ..
ENSG00000188976 1.76549e-35      NOC2L     26155 NOC2 like nucleolar ..
ENSG00000187961 1.13413e-07      KLHL17    339451 kelch like family me..
ENSG00000187583 9.19031e-01      PLEKHN1   84069 pleckstrin homology ..
ENSG00000187642 4.03379e-01      PERM1     84808 PPARGC1 and ESRR ind..
ENSG00000188290 1.30538e-24      HES4      57801 hes family bHLH tran..
ENSG00000187608 2.37452e-02      ISG15     9636 ISG15 ubiquitin like..
ENSG00000188157 4.21963e-16      AGRN      375790          agrin
ENSG00000237330          NA      RNF223    401934 ring finger protein ..
```

Q. Finally for this section let's reorder these results by adjusted p-value and save them to a CSV file in your current project directory.

```
res = res[order(res$pvalue),]
write.csv(res, file ="deseq_results.csv")
```

## (2) Pathway Analysis

Now we can load the packages and setup the KEGG data-sets we need.

```

library(pathview)
library(gage)
library(gageData)

data(kegg.sets.hs)
data(sigmet.idx.hs)

# Focus on signaling and metabolic pathways only
kegg.sets.hs = kegg.sets.hs[sigmet.idx.hs]

# Examine the first 3 pathways
head(kegg.sets.hs, 3)

$`hsa00232 Caffeine metabolism`
[1] "10"    "1544"  "1548"  "1549"  "1553"  "7498"  "9"

$`hsa00983 Drug metabolism - other enzymes`
[1] "10"    "1066"  "10720" "10941" "151531" "1548"  "1549"  "1551"
[9] "1553"  "1576"  "1577"  "1806"  "1807"   "1890"  "221223" "2990"
[17] "3251"  "3614"  "3615"  "3704"  "51733"  "54490" "54575"  "54576"
[25] "54577" "54578" "54579" "54600" "54657"  "54658" "54659"  "54963"
[33] "574537" "64816" "7083"  "7084"  "7172"   "7363"  "7364"   "7365"
[41] "7366"  "7367"  "7371"  "7372"  "7378"   "7498"  "79799" "83549"
[49] "8824"  "8833"  "9"     "978"

$`hsa00230 Purine metabolism`
[1] "100"   "10201" "10606" "10621" "10622" "10623" "107"   "10714"
[9] "108"   "10846" "109"   "111"   "11128" "11164" "112"   "113"
[17] "114"   "115"   "122481" "122622" "124583" "132"   "158"   "159"
[25] "1633"  "171568" "1716"  "196883" "203"   "204"   "205"   "221823"
[33] "2272"  "22978" "23649" "246721" "25885" "2618"  "26289" "270"
[41] "271"   "27115" "272"   "2766"  "2977"  "2982"  "2983"  "2984"
[49] "2986"  "2987"  "29922" "3000"  "30833" "30834" "318"   "3251"
[57] "353"   "3614"  "3615"  "3704"  "377841" "471"   "4830"  "4831"
[65] "4832"  "4833"  "4860"  "4881"  "4882"  "4907"  "50484" "50940"
[73] "51082" "51251" "51292" "5136"  "5137"  "5138"  "5139"  "5140"
[81] "5141"  "5142"  "5143"  "5144"  "5145"  "5146"  "5147"  "5148"
[89] "5149"  "5150"  "5151"  "5152"  "5153"  "5158"  "5167"  "5169"
[97] "51728" "5198"  "5236"  "5313"  "5315"  "53343" "54107" "5422"
[105] "5424"  "5425"  "5426"  "5427"  "5430"  "5431"  "5432"  "5433"
[113] "5434"  "5435"  "5436"  "5437"  "5438"  "5439"  "5440"  "5441"

```

```
[121] "5471"    "548644"   "55276"    "5557"     "5558"     "55703"    "55811"    "55821"
[129] "5631"    "5634"     "56655"    "56953"    "56985"    "57804"    "58497"    "6240"
[137] "6241"    "64425"    "646625"   "654364"   "661"      "7498"     "8382"     "84172"
[145] "84265"   "84284"    "84618"    "8622"     "8654"     "87178"    "8833"     "9060"
[153] "9061"    "93034"    "953"      "9533"     "954"      "955"      "956"      "957"
[161] "9583"    "9615"
```

The main gage() function requires a named vector of fold changes, where the names of the values are the Entrez gene IDs.

Note that we used the mapIDs() function above to obtain Entrez gene IDs (stored in `res$entrez`) and we have the fold change results from DESeq2 analysis (stored in `res$log2FoldChange`).

```
foldchanges = res$log2FoldChange
names(foldchanges) = res$entrez
head(foldchanges)
```

```
1266      54855      1465      51232      2034      2317
-2.422719  3.201955 -2.313738 -2.059631 -1.888019 -1.649792
```

Now, let's run the gage pathway analysis.

```
# Get the results
keggres = gage(foldchanges, gsets=kegg.sets.hs)
```

Now lets look at the object returned from gage().

```
attributes(keggres)
```

```
$names
[1] "greater" "less"    "stats"
```

It is a list with three elements, “greater”, “less” and “stats”.

You can also see this in your Environment panel/tab window of RStudio or use the R command `str(keggres)`.

Like any list we can use the dollar syntax to access a named element, e.g. `head(keggres$greater)` and `head(keggres$less)`.

Lets look at the first few down (less) pathway results:

```
# Look at the first few down (less) pathways
head(keggres$less)
```

	p.geomean	stat.mean	p.val
hsa04110 Cell cycle	8.995727e-06	-4.378644	8.995727e-06
hsa03030 DNA replication	9.424076e-05	-3.951803	9.424076e-05
hsa03013 RNA transport	1.375901e-03	-3.028500	1.375901e-03
hsa03440 Homologous recombination	3.066756e-03	-2.852899	3.066756e-03
hsa04114 Oocyte meiosis	3.784520e-03	-2.698128	3.784520e-03
hsa00010 Glycolysis / Gluconeogenesis	8.961413e-03	-2.405398	8.961413e-03
	q.val	set.size	exp1
hsa04110 Cell cycle	0.001448312	121	8.995727e-06
hsa03030 DNA replication	0.007586381	36	9.424076e-05
hsa03013 RNA transport	0.073840037	144	1.375901e-03
hsa03440 Homologous recombination	0.121861535	28	3.066756e-03
hsa04114 Oocyte meiosis	0.121861535	102	3.784520e-03
hsa00010 Glycolysis / Gluconeogenesis	0.212222694	53	8.961413e-03

Each keggres\$less and keggres\$greater object is data matrix with gene sets as rows sorted by p-value.

The top “less/down” pathways is “Cell cycle” with the KEGG pathway identifier hsa04110.


Now, let’s try out the pathview() function from the pathview package to make a pathway plot with our RNA-Seq expression results shown in color. To begin with lets manually supply a pathway.id (namely the first part of the “hsa04110 Cell cycle”) that we could see from the print out above.

```
pathview(gene.data=foldchanges, pathway.id="hsa04110")
```

```
'select()' returned 1:1 mapping between keys and columns
```

```
Info: Working in directory C:/Users/ambro/OneDrive/Desktop/BIMM 143/Class14
```

```
Info: Writing image file hsa04110.pathview.png
```



Data on KEGG graph  
Rendered by Pathview

```
pathview(gene.data=foldchanges,  
pathway.id="hsa04110", kegg.native=FALSE) Now, let's process our results a bit more to  
automagically pull out the top 5 upregulated pathways, then further process that just to get  
the pathway IDs needed by the pathview() function. We'll use these KEGG pathway IDs for  
pathview plotting below.
```

```
## Focus on top 5 upregulated pathways here for demo purposes only
keggrespathways <- rownames(keggres$greater)[1:5]
```

```
# Extract the 8 character long IDs part of each string  
keggresids = substr(keggrespathways, start=1, stop=8)  
keggresids
```

```
[1] "hsa04640" "hsa04630" "hsa00140" "hsa04142" "hsa04330"
```

Finally, lets pass these IDs in keggresids to the pathview() function to draw plots for all the top 5 pathways.

```
pathview(gene.data=foldchanges, pathway.id=keggresids, species="hsa")
```

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory C:/Users/ambro/OneDrive/Desktop/BIMM 143/Class14

Info: Writing image file hsa04640.pathview.png

```
'select()' returned 1:1 mapping between keys and columns

Info: Working in directory C:/Users/ambro/OneDrive/Desktop/BIMM 143/Class14

Info: Writing image file hsa04630.pathview.png

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory C:/Users/ambro/OneDrive/Desktop/BIMM 143/Class14

Info: Writing image file hsa00140.pathview.png

'select()' returned 1:1 mapping between keys and columns



Info: Working in directory C:/Users/ambro/OneDrive/Desktop/BIMM 143/Class14


Info: Writing image file hsa04142.pathview.png

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory C:/Users/ambro/OneDrive/Desktop/BIMM 143/Class14

Info: Writing image file hsa04330.pathview.png
```





Q. Can you do the same procedure as above to plot the pathview figures for the top 5 down-regulated pathways?

```

## Focus on top 5 upregulated pathways here for demo purposes only
keggrespathways.down <- rownames(keggres$less)[1:5]

# Extract the 8 character long IDs part of each string
keggresids.down = substr(keggrespathways.down, start=1, stop=8)
keggresids.down

[1] "hsa04110" "hsa03030" "hsa03013" "hsa03440" "hsa04114"

pathview(gene.data=foldchanges, pathway.id=keggresids.down, species="hsa")

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory C:/Users/ambro/OneDrive/Desktop/BIMM 143/Class14

Info: Writing image file hsa04110.pathview.png

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory C:/Users/ambro/OneDrive/Desktop/BIMM 143/Class14

Info: Writing image file hsa03030.pathview.png

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory C:/Users/ambro/OneDrive/Desktop/BIMM 143/Class14

Info: Writing image file hsa03013.pathview.png

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory C:/Users/ambro/OneDrive/Desktop/BIMM 143/Class14


Info: Writing image file hsa03440.pathview.png


'select()' returned 1:1 mapping between keys and columns

```

Info: Working in directory C:/Users/ambro/OneDrive/Desktop/BIMM 143/Class14

Info: Writing image file hsa04114.pathview.png





### (3) Gene Ontology (GO)

We can also do a similar procedure with gene ontology. Similar to above, go.sets.hs has all GO terms. go.subs.hs is a named list containing indexes for the BP, CC, and MF ontologies. Let's focus on BP (a.k.a Biological Process) here.

```

data(go.sets.hs)
data(go.subs.hs)

# Focus on Biological Process subset of GO
gobpsets = go.sets.hs[go.subs.hs$BP]

gobpres = gage(foldchanges, gsets=gobpsets)

lapply(gobpres, head)

```

\$greater

	p.geomean	stat.mean	p.val
GO:0007156 homophilic cell adhesion	8.519724e-05	3.824205	8.519724e-05
GO:0002009 morphogenesis of an epithelium	1.396681e-04	3.653886	1.396681e-04
GO:0048729 tissue morphogenesis	1.432451e-04	3.643242	1.432451e-04
GO:0007610 behavior	1.925222e-04	3.565432	1.925222e-04
GO:0060562 epithelial tube morphogenesis	5.932837e-04	3.261376	5.932837e-04
GO:0035295 tube development	5.953254e-04	3.253665	5.953254e-04
	q.val	set.size	exp1
GO:0007156 homophilic cell adhesion	0.1951953	113	8.519724e-05
GO:0002009 morphogenesis of an epithelium	0.1951953	339	1.396681e-04
GO:0048729 tissue morphogenesis	0.1951953	424	1.432451e-04
GO:0007610 behavior	0.1967577	426	1.925222e-04
GO:0060562 epithelial tube morphogenesis	0.3565320	257	5.932837e-04
GO:0035295 tube development	0.3565320	391	5.953254e-04

\$less

	p.geomean	stat.mean	p.val
GO:0048285 organelle fission	1.536227e-15	-8.063910	1.536227e-15
GO:0000280 nuclear division	4.286961e-15	-7.939217	4.286961e-15
GO:0007067 mitosis	4.286961e-15	-7.939217	4.286961e-15
GO:0000087 M phase of mitotic cell cycle	1.169934e-14	-7.797496	1.169934e-14
GO:0007059 chromosome segregation	2.028624e-11	-6.878340	2.028624e-11
GO:0000236 mitotic prometaphase	1.729553e-10	-6.695966	1.729553e-10
	q.val	set.size	exp1
GO:0048285 organelle fission	5.841698e-12	376	1.536227e-15
GO:0000280 nuclear division	5.841698e-12	352	4.286961e-15
GO:0007067 mitosis	5.841698e-12	352	4.286961e-15
GO:0000087 M phase of mitotic cell cycle	1.195672e-11	362	1.169934e-14
GO:0007059 chromosome segregation	1.658603e-08	142	2.028624e-11
GO:0000236 mitotic prometaphase	1.178402e-07	84	1.729553e-10

```
$stats
stat.mean      exp1
GO:0007156 homophilic cell adhesion      3.824205 3.824205
GO:0002009 morphogenesis of an epithelium 3.653886 3.653886
GO:0048729 tissue morphogenesis          3.643242 3.643242
GO:0007610 behavior                      3.565432 3.565432
GO:0060562 epithelial tube morphogenesis 3.261376 3.261376
GO:0035295 tube development              3.253665 3.253665
```

## (5) Reactome Analysis

Reactome is database consisting of biological molecules and their relation to pathways and processes. Reactome, such as many other tools, has an online software available <https://reactome.org/> and R package available <https://bioconductor.org/packages/release/bioc/html/ReactomePA.html>

```
sig_genes <- res[res$padj <= 0.05 & !is.na(res$padj), "symbol"]
print(paste("Total number of significant genes:", length(sig_genes)))
```

```
[1] "Total number of significant genes: 8147"
```

```
write.table(sig_genes, file="significant_genes.txt", row.names=FALSE, col.names=FALSE, quote=
```

Q: What pathway has the most significant “Entities p-value”? Do the most significant pathways listed match your previous KEGG results? What factors could cause differences between the two methods?

The mitotic cell pathway has the most significance, with a p-value of 2.02E-5. They do match. The cell cycle was the most significant in the KEGG analysis. The two methods operate at different levels of resolution, with KEGG identifying the broader cell cycle pathway and the other analysis specifying the mitotic phase as a more detailed subset.